

Glucocorticoid-Responsive Genes in Airway Smooth Muscle Cells: A Comparison of Modern and the Original RNA-seq Pipelines

Ilia Gogotchuri

This is a replication study of the 2014 paper by Blanca E. Himes and Xiaofeng Jiang [1]

Abstract.

RNA-seq analysis pipelines have evolved significantly since 2014, with improvements in alignment algorithms, quantification methods, and analysis tools. Using a newer pipeline and reference genome, we reanalyzed the GSE52778 RNA-seq dataset, which contains data from eight airway smooth muscle cells (ASM) from four donors. The dataset was used in the original paper to characterize transcriptomic changes in human ASM cells treated with dexamethasone (potent glucocorticoid). The original paper used the Tuxedo tools, TopHat2 for alignment, and Cufflinks+Cuffdiff for quantification and differential expression analysis, with the reference genome **hg19**. In this study, we modernized the pipeline using STAR, featureCounts, and DESeq2, and also compared the reference-genome effect by running the pipeline on both the original reference genome **hg19** and the latest available reference genome **hg38**. The STAR alignment achieved a **97.6%** mapping rate compared to **83.4%** in the original paper. Junction-spanning reads were also improved from 26% to 45%. Despite the technical metric improvements, we have rediscovered and **validated 96.2%** (**304/316**) of the genes identified in the original paper (with **hg19**), with a Pearson correlation of 0.998 between the log₂ fold changes. The modern pipeline with the new reference genome also **identified 909** significant differentially expressed genes ($padj < 0.05$ and $|log2FoldChange| > 1$), reflecting higher sensitivity. Although the reference genome effect was not significant, it might still be interesting for biological interpretation. We conclude that the original well-designed RNA-seq analysis was robust to pipeline variations and would likely still benefit from the modern tooling.

Key words: RNA-seq; bioinformatics; differential expression; STAR alignment; DESeq2; glucocorticoid response

1. Introduction

1.1. Pipeline Background

RNA-seq is a widely used method to study transcriptomes in cells. The alignment tools evolved significantly in the last decade. One of the most effective tools for alignment we have today—STAR, was first introduced in a paper published at the end of 2012 [2]. In 2013, a paper [3] was published comparing the performance of several alignment tools. GSNAp, GSTRUCT, MapSplice, and STAR emerged as the top performers. But the interesting thing about STAR was its performance: it was about 180× faster than GSNAp and MapSplice while maintaining comparable accuracy and, in some cases, improving junction-spanning with two passes. This allowed STAR to run on a single modern computer and perform alignment in a few hours instead of days, making discovery and iteration faster, cheaper, and overall more accessible. Running the scripts in this repository took around 2 hours in total.

Table 1: Samples Table

Sample ID	SRR ID	Condition	Cell Line
N61311_untreated	SRR1039508	untreated	N61311
N61311_dex	SRR1039509	dex	N61311
N052611_untreated	SRR1039512	untreated	N052611
N052611_dex	SRR1039513	dex	N052611
N080611_untreated	SRR1039516	untreated	N080611
N080611_dex	SRR1039517	dex	N080611
N061011_untreated	SRR1039520	untreated	N061011
N061011_dex	SRR1039521	dex	N061011

For assigning sequence reads to genomic features, the pipeline uses featureCounts [4], another highly efficient tool that is an order of magnitude faster than comparable tools like HTSeq. For fold change and dispersion analysis, the pipeline uses a Python version of DESeq2 [5].

1.2. Dataset

The dataset used in the paper is available at GSE52778 (<https://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE52778>). The dataset actually has 16 samples, combinations of untreated, treated with Albuterol, and treated with Dexamethasone. Since the study mostly focuses on untreated and dexamethasone-treated (12 hours) samples, we will use only these two groups. The samples are listed in Table 1: four donors and two conditions.

1.3. Original Study Summary

The original paper used Tuxedo tools: TopHat2 for alignment and Cufflinks+Cuffdiff for quantification and differential expression analysis, with the reference genome **hg19**. Analysis of the original article identified 316 genes that were significantly differentially expressed between the two conditions ($p < 0.05$). The **CRISPLD2** gene was identified as a novel finding by the paper. The top gene findings of the original paper are shown in Table 2 (padj $< 1 \times 10^{-16}$ is marked 0).

1.4. Study Design and Objectives

You can find the original analysis description in the original paper, on page 9, paragraph “RNA-Seq Data Analysis.”

This study tries to separate **pipeline effects** from **reference genome effects** by performing two-part analyses:

We will try to validate the findings from the original paper and compare them to the findings from the modern pipeline.

1.4.1 Expected Insights

- **Part 1 vs Original:** Effect of the updated pipeline
- **Part 2 vs Part 1:** Effect of updated genome assembly and annotation
- **Part 2 vs Original:** Combined modernization effect

Table 2: Top DEGs from the original paper

Gene	log2 FC	P adj (Q)
C7	-3.35	0
CCDC69	-2.92	0
DUSP1	-2.99	0
FKBP5	-3.95	0
GPX3	-3.76	0
KLF15	-4.58	0
MAOA	-3.29	0
SAMHD1	-3.83	0
SERPINA3	-3.34	0
SPARCL1	-4.70	0
C13orf15	-3.27	2.5×10^{-13}
TSC22D3	-3.27	2.5×10^{-13}
CRISPLD2	-2.70	6.9×10^{-13}

Table 3: Pipeline Effects

Analysis	Reference	Pipeline	Purpose
Original	hg19	TopHat/Cufflinks + Cuffdiff	Baseline (2014)
Part 1	hg19	STAR/featureCount + DESeq2	Isolate pipeline improvement
Part 2	hg38	STAR/featureCount + DESeq2	Full modernization

Table 4: Alignment Summary

Metric	hg19 (TopHat2)	hg19 (STAR+fc)	hg38 (STAR+fc)
Reads (avg)	58.9M	48.9M	48.9M
Reads (min)	44.2M	33.7M	39.7M
Reads (max)	71.3M	68.6M	68.6M
Mapped (avg)	83.36%	97.59%	97.64%
Mapped (min)	81.94%	96.26%	96.33%
Mapped (max)	84.34%	98.10%	98.13%
Junctions (avg)	26.43%	45.57%	44.52%
Junctions (min)	NA	43.06%	42.07%
Junctions (max)	NA	47.10%	45.94%

2. Results

2.1. Alignment Quality

Modern STAR alignment achieved a 97.6% mapping rate compared to 83.4% in the original paper using TopHat2. Junction-spanning reads were also improved from 26% to 45%, indicating improvement in splice site detection by STAR. Although the read counts are lower.

Table 5: DEG Summary

Metric	Original	hg19 Modern	hg38 Modern
Genes Tested	20,561	20,561	22,136
DEGs ($p_{adj} < 0.05$)	316	2,648	2,735
Significant DEGs ($p_{adj} < 0.05$, $ LFC > 1$)	N/A	872	909
Upregulated	99	396	420
Downregulated	219	476	489




Figure 1: Volcano plots show a similar distribution of gene expression changes across the three genome analyses. Genes are predominantly downregulated in all of them.

2.2. Gene Differential Expression Analysis

The original paper identified 316 genes significantly differentially expressed between the two conditions with $p < 0.05$. The modern pipeline with the new reference genome also identified a total of 910 significant differentially expressed genes ($p_{adj} < 0.05$ and $|log2FoldChange| > 1$), reflecting higher sensitivity of DESeq2 analysis.

2.3. DEG Correlations and Validation

We have re-discovered and validated 96.2% (304/316) (with **hg19**) of the genes identified by the original paper with 0.998 Pearson correlation between the log2 fold changes. With the updated genome **hg38**, we have validated 92.4% (292/316) of the genes with the correlation remaining the same at 0.998. Confirming both methods detect similar biological effects.

Although the validation rate drop in the modern genome is mostly due to gene symbol changes, annotation updates, or minor coordinate shifts between genome builds rather than loss of biological signal. I have manually tracked down some (7) of the gene symbol changes, the figures do not reflect the corrections, but I have adjusted the percentage of validated genes accordingly.

2.4. Novel Findings

After analyzing the modern genome **hg38**, we have identified several significant DEGs that were not reported in the original paper 6. Those might not have been detectable with the older




Figure 2: Analysis with the updated genome validated one additional significant DEG (compared to hg19) from the original findings—**SERPINA3** and didn’t find as strong of an effect with 16 genes.

Table 6: Novel DEGs identified with HG38 (10 novel genes identified by the modern pipeline with hg38 reference genome having padj (Q value) $< 1 \times 10^{-10}$ and $|\log_{2}\text{FC}| > 1$)

Gene	padj	log2FC
GASK1B (FAM198B)	6.87×10^{-60}	-1.12
RAB7B	4.66×10^{-26}	1.83
ENSG00000250978	8.37×10^{-19}	-6.30
MARCHF10	1.87×10^{-18}	-3.96
LINC02884	2.10×10^{-13}	-3.13
RELL1	1.52×10^{-11}	-1.12
HMGA2-AS1	3.61×10^{-11}	-2.35
PEAK1	8.70×10^{-11}	-1.00
MIR99AHG	2.45×10^{-10}	1.51
AOPEP	8.45×10^{-10}	-1.10

pipeline or genome version. And those might be interesting for further study and biological interpretation.




Figure 3: Between the original and modern analyses, the log2 fold changes are almost perfectly correlated with a Pearson correlation of 0.998 for both genome versions (regardless of the reference genome used for alignment).

3. Discussion

3.1. Modernization of the Alignment and Analysis Pipeline

The transition from TopHat2/Cufflinks to STAR/featureCounts/DESeq2 resulted in technical improvements while preserving biological findings. Alignment rates increased by 14 percentage points ($83.4\% \rightarrow 97.6\%$), indicating that STAR’s splice site detection algorithm works better to map RNA-seq reads to the reference genome. Junction-spanning reads were also improved from 26% to 45%, indicating that STAR’s junction-spanning detects splice sites more accurately.

Despite the improved sensitivity and higher alignment rates, the biological findings remained remarkably consistent. The near-perfect Pearson correlation and 96.2% gene rediscovery rate likely indicate that the transcriptomic response in ASM to dexamethasone treatment is reproducible and robust to methodological variation. The key Glucocorticoid-response genes highlighted in the original paper were also identified in the modern pipeline, with near-perfect




Figure 4: Looking at the top DEGs and their LFCs, we can see that they are almost identical between the modern and the original analyses. The DEGs with their directions and magnitudes of changes are also identical.

correlation.

Hence, I can assume that the increased number of DEGs reflects improved alignment rates and a more sensitive negative binomial analysis model in DESeq2, rather than false positives.

3.2. New Reference Genome Effects

Comparison between hg19 and hg38 analyses revealed minimal impact from reference genome updates on differential expression results. However, the updated genome did enable identifi-

cation of 10 novel DEGs with high confidence ($\text{padj} < 1 \times 10^{-10}$, $|\log_{2}\text{FC}| > 1$) that were not detectable with hg19. These include genes like GASK1B, RAB7B and MARCHF10, which may represent improved annotations in GRCh38 or sequence corrections that allow more accurate read mapping.

The SERPINA3 gene, validated in hg38 but not in hg19, is an interesting case. It was also discovered with hg19, but the signal was not as strong. Why the newer reference genome alignment allowed us to detect it is still an open question. The slight decrease in validated original genes (96.2% → 92.4%) when moving to hg38 likely reflects gene symbol changes, annotation updates, or minor coordinate shifts between genome builds rather than loss of biological signal. Researchers should consider that direct gene-level comparisons across genome versions may require additional symbol mapping.

3.3. Novel Findings

The novel findings, especially ENSG00000250978 and MARCHF10 seem to have a strong response to dexamethasone treatment, and their Q values are very low, statistically unlikely to be false positives. Many of the new findings are not well-studied genes, especially in the context of glucocorticoid response, and there is hardly any literature on the two I have highlighted.

4. Conclusion

This reanalysis demonstrates that the original study produced robust findings that withstand methodological variations. Using STAR alignment and DESeq2 differential expression analysis, we validated 96.2% of the original 316 differentially expressed genes with near-perfect fold change correlation ($r = 0.998$).

Modern tools provided significant improvements: increased alignment rates from 83% to 97.6% and the detection of splice junctions has nearly doubled. These improvements translated to increased statistical sensitivity, which allowed us to identify a total of 909 DEGs with high confidence ($\text{padj} < 0.05$, $|\log_{2}\text{FC}| > 1$). Which, compared to the original paper, is nearly threefold increase, while maintaining high specificity, suggested by almost perfect correlation of the validated genes.

Reference genome version update from **hg19** to **hg38** had minimal impact on core findings, but enabled the discovery of 10 novel candidate genes that might be of interest for further study.

5. Methods: Pipeline Setup and Project Structure

5.1. Environment Setup

Github Repository contains all the necessary components for replicating this study

All analysis was performed using conda environment management. The `environment.yml` file contains all dependency declarations:

```
conda env create -f environment.yml
conda activate bioinf-grg
```

5.2. Data Gathering Automation

Data gathering is automated with `scripts/01_download_data.py`. Simply running the script will download all the data required for the analysis. Configuration for the data to be used

is in `config/samples_table.tsv` and `config/configuration.py`. Reference genomes are downloaded using `wget` from the datasets of EMBL's European Bioinformatics Institute and placed under `data/reference/{hg19|hg38}`. Fasta files are downloaded using `fasterq-dump` from SRA and for compression using `pigz`. Fasta files are placed under `data/raw/`.

Additionally, I have manually downloaded and formatted tables from the original paper supplementary materials. The 316 genes identified by the original paper are from the supplementary table S3. I modified the headers to correspond to the correct columns produced by the newer pipeline. All files are placed under `config/`.

5.3. Alignment

Part A: STAR index is built using `scripts/02a_build_star_index.py`. This script takes around an hour to run for the two genomes and needs \geq 32GB of RAM.

Part B: Alignment is performed using `scripts/02b_align_reads.py`. This uses STAR index built in Part A and takes around 5 minutes per sample, producing BAM files for each sample. After this part, every next step is significantly faster.

5.4. Gene Quantification

`featureCounts` uses the produced BAM files to quantify gene expression. `featureCounts` is run using `scripts/03_quantify_genes.py`. The script contains post-processing and cleaning steps to produce the final table under `results/hg{19,38}/tables/gene_counts.tsv` containing gene counts for each sample.

5.5. Differential Expression Analysis

Differential expression analysis is performed using `scripts/04_analysis_deseq2.py`. This script uses the gene counts table produced in the previous step and performs DESeq2 analysis. The complete list of results with $\text{padj} < 0.05$ is saved under `results/hg{19,38}/tables/deseq2_results.csv`. The significant genes with $|\log2\text{FoldChange}| > 1$ are saved under `results/hg{19,38}/tables/top_results.csv`.

5.6. Comparisons and Visualization

The results from Part 1 and Part 2 are compared to the original findings and to each other using scripts `05a_comparisons_to_original.py` and `05b_comparisons_between_genomes.py`. General statistics are collected using `06_collect_general_stats.py`. All figures are produced using `07_plots.py`.

References

- [1] Himes BE, Jiang X, Wagner P, Hu R, Wang Q, et al. (2014) RNA-Seq Transcriptome Profiling Identifies CRISPLD2 as a Glucocorticoid Responsive Gene that Modulates Cytokine Function in Airway Smooth Muscle Cells. PLOS ONE 9(6): e99625. <https://doi.org/10.1371/journal.pone.0099625>
- [2] Dobin A, Davis CA, Schlesinger F, et al. STAR: ultrafast universal RNA-seq aligner. Bioinformatics. 2013;29(1):15-21. doi:10.1093/bioinformatics/bts635
- [3] Engström, P., Steijger, T., Sipos, B. et al. Systematic evaluation of spliced alignment programs for RNA-seq data. Nat Methods 10, 1185–1191 (2013). <https://doi.org/10.1038/nmeth.2722>

- [4] Yang Liao, Gordon K. Smyth, Wei Shi, featureCounts: an efficient general purpose program for assigning sequence reads to genomic features, *Bioinformatics*, Volume 30, Issue 7, April 2014, Pages 923–930, <https://doi.org/10.1093/bioinformatics/btt656>
- [5] Love, M. I., Huber, W., Anders, S. (2014). Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2. *Genome biology*, 15(12), 550. <https://doi.org/10.1186/s13059-014-0550-8>